

AI Assisted Coding

Assignment – 17.5

Name: K. Sai Krishna Reddy

RollNo: 2303A52305

Batch: 45

Task 1 –Loan Approval Dataset Cleaning

Use AI to generate preprocessing code for a loan dataset.

Instructions:

- Handle missing values in loan_amount, credit_score.
- Encode loan_status (Approved/Rejected).
- Scale numerical features using standard scaling.
- Split the dataset into training and testing sets.

Sample Code:

```
import pandas as pd  
  
data = pd.read_csv("loadapproval.csv")  
  
print(data.head())
```

Dataset link:

<https://www.kaggle.com/datasets/amineipad/loan-approval-dataset>

Expected Output:

- A machine-learning-ready loan dataset.

```
import pandas as pd
data = pd.read_csv('/content/drive/MyDrive/loanapproval.csv')
print(data.head())
```

	applicant_id	age	gender	marital_status	annual_income	loan_amount	
0	1	59	Male	Divorced	100073	7169	
1	2	49	Male	Married	112197	23556	
2	3	35	Male	Divorced	84429	27052	
3	4	63	Female	Single	124195	11313	
4	5	28	Female	Married	81627	13315	

	credit_score	num_dependents	existing_loans_count	employment_status	
0	793	1	1	Unemployed	
1	789	0	2	Employed	
2	372	1	4	Unemployed	
3	808	3	4	Self-employed	
4	689	0	1	Unemployed	

	loan_approved
0	1
1	1
2	0
3	1
4	1

```
# Check for missing values in 'loan_amount' and 'credit_score'
print("Missing values before imputation:")
print(data[['loan_amount', 'credit_score']].isnull().sum())

# Impute missing values with the median for 'loan_amount' and 'credit_score'
# Using .copy() to avoid SettingWithCopyWarning, ensuring changes are made on a copy of the slice.
data_imputed = data.copy()
# Updated to avoid FutureWarning by reassigning the column instead of using inplace=True
data_imputed['loan_amount'] = data_imputed['loan_amount'].fillna(data_imputed['loan_amount'].median())
data_imputed['credit_score'] = data_imputed['credit_score'].fillna(data_imputed['credit_score'].median())

print("\nMissing values after imputation:")
print(data_imputed[['loan_amount', 'credit_score']].isnull().sum())

# Display the first few rows of the updated DataFrame to show the effect of imputation
print("\nDataFrame head after imputation:")
display(data_imputed.head())
```

Missing values before imputation:

```
loan_amount    0
credit_score    0
dtype: int64
```

Missing values after imputation:

```
loan_amount    0
credit_score    0
dtype: int64
```

DataFrame head after imputation:

	applicant_id	age	gender	marital_status	annual_income	loan_amount	credit_score	num_dependents	existing_loans_count
0	1	59	Male	Divorced	100073	7169	793	1	1
1	2	49	Male	Married	112197	23556	789	0	2
2	3	35	Male	Divorced	84429	27052	372	1	4
3	4	63	Female	Single	124195	11313	808	3	4
4	5	28	Female	Married	81627	13315	689	0	1

Handle missing values in loan_amount, credit_score.

I will check for missing values in the 'loan_amount' and 'credit_score' columns and then impute them using the median. This approach is generally good for numerical data as it's less sensitive to outliers compared to the mean.

The analysis shows that there were no missing values in the 'loan_amount' and 'credit_score' columns, so the imputation step didn't change the data. This is great, as it means your data is already clean for these columns! You also received a `FutureWarning` from pandas about using `inplace=True` with `fillna` on a copied DataFrame. While it didn't cause an error this time, I'll update the code to use a more robust assignment method (`data_imputed['column'] = data_imputed['column'].fillna(...)`) to prevent potential issues in future pandas versions.

The code executed successfully! It shows that there were no missing values in the 'loan_amount' or 'credit_score' columns, both before and after the imputation step. This means your dataset already has complete information for these features. The `FutureWarning` that appeared in the previous run has also been addressed, ensuring the code is up-to-date with pandas best practices.

Calculate descriptive statistics for the 'loan_amount' and 'credit_score' columns
Visualize the distribution of 'loan_amount' using a histogram
Analyze the correlation between 'credit_score' and 'loan_approved' status

What can I help you build?

Gemini 2.5 Flash

```
# Map 0 to 'Rejected' and 1 to 'Approved' in the 'loan_approved' column
# We will apply this transformation on the data_imputed DataFrame.
loan_status_mapping = {0: 'Rejected', 1: 'Approved'}
data_imputed['loan_status'] = data_imputed['loan_approved'].map(loan_status_mapping)

# Display the updated DataFrame with the new 'loan_status' column
print("DataFrame head with encoded loan status:")
display(data_imputed[['loan_approved', 'loan_status']].head())
```

DataFrame head with encoded loan status:

	loan_approved	loan_status
0	1	Approved
1	1	Approved
2	0	Rejected
3	1	Approved
4	1	Approved

The code executed successfully! It shows that there were no missing values in the 'loan_amount' or 'credit_score' columns, both before and after the imputation step. This means your dataset already has complete information for these features. The `FutureWarning` that appeared in the previous run has also been addressed, ensuring the code is up-to-date with pandas best practices.

Encode loan_status (Approved/Rejected).

I will encode the 'loan_approved' column by mapping 1 to 'Approved' and 0 to 'Rejected'. This will convert the numerical status into more descriptive labels.

The 'loan_approved' column has been successfully encoded into a new 'loan_status' column. As you can see from the output, a value of 1 in 'loan_approved' now corresponds to 'Approved' in 'loan_status', and 0 corresponds to 'Rejected'. This makes the loan approval status much more readable.

```
from sklearn.preprocessing import StandardScaler

# Identify numerical features to be scaled (excluding applicant_id, loan_approved, and loan_status)
numerical_features = ['age', 'annual_income', 'loan_amount', 'credit_score', 'num_dependents', 'existing_loans_count']

# Initialize the StandardScaler
scaler = StandardScaler()

# Fit and transform the numerical features
data_scaled_numerical = scaler.fit_transform(data_imputed[numerical_features])

# Create a DataFrame from the scaled numerical features
data_scaled = pd.DataFrame(data_scaled_numerical, columns=numerical_features)

# Optionally, you can add back the non-numerical columns to the scaled DataFrame
# For now, let's just display the scaled numerical features

print("Head of DataFrame with scaled numerical features:")
display(data_scaled.head())
```

Head of DataFrame with scaled numerical features:

	age	annual_income	loan_amount	credit_score	num_dependents	existing_loans_count
0	1.307840	0.482301	-1.566427	1.391620	-0.689042	-0.737500
1	0.514489	0.805362	-0.287825	1.365954	-1.394305	-0.029726
2	-0.596204	0.065444	-0.015048	-1.309808	-0.689042	1.385820
3	1.625181	1.125066	-1.243090	1.487871	0.721484	1.385820
4	-1.151550	-0.009219	-1.086883	0.724284	-1.394305	-0.737500

The 'loan_approved' column has been successfully encoded into a new 'loan_status' column. As you can see from the output, a value of 1 in 'loan_approved' now corresponds to 'Approved' in 'loan_status', and 0 corresponds to 'Rejected'. This makes the loan approval status much more readable.

Scale numerical features using standard scaling.

I'll scale the numerical features in your dataset using `StandardScaler`. This will transform the data so that it has a mean of 0 and a standard deviation of 1, which is often beneficial for many machine learning algorithms. I'll display the first few rows of the scaled data afterwards.

The numerical features (age, annual income, loan amount, credit score, num dependents, and existing loans count) have been successfully scaled using `StandardScaler`. This transformation normalizes the data to have a mean of 0 and a standard deviation of 1, which is crucial for many machine learning algorithms. The output above shows the first few rows of the DataFrame with these scaled numerical features.

Visualize the distribution of 'annual_income' using a histogram
Perform one-hot encoding on categorical features like 'gender' and 'employment_status'
Split the dataset into training and testing sets for model development

```

from sklearn.model_selection import train_test_split

# Define features (X) and target (y)
# Drop 'applicant_id' as it's an identifier and not a predictive feature.
# Drop 'loan_approved' from X as it's the target variable.
# Drop 'loan_status' as it's an encoded version of the target for display.
X = data_imputed.drop(columns=['applicant_id', 'loan_approved', 'loan_status'])
y = data_imputed['loan_approved']

# Perform one-hot encoding for categorical features in X
# Identify categorical columns
categorical_cols = X.select_dtypes(include=['object']).columns
X_encoded = pd.get_dummies(X, columns=categorical_cols, drop_first=True) # drop_first avoids multicollinearity

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.2, random_state=42)

print("Shapes of the split datasets:")
print(f"X_train shape: {X_train.shape}")
print(f"X_test shape: {X_test.shape}")
print(f"y_train shape: {y_train.shape}")
print(f"y_test shape: {y_test.shape}")

# Display first few rows of X_train and y_train
print("\nX_train head:")
display(X_train.head())
print("\ny_train head:")
display(y_train.head())

Shapes of the split datasets:
X_train shape: (800, 11)
X_test shape: (200, 11)
y_train shape: (800,)
y_test shape: (200,)

X_train head:

```

	age	annual_income	loan_amount	credit_score	num_dependents	existing_loans_count	gender_Male	marital_status_Married	marit
29	47	83646	26300	428	1	4	True	False	
535	56	92954	26542	301	4	4	False	True	
695	48	134604	35581	329	2	3	True	True	
557	59	28472	48463	397	4	4	True	False	

Gemini

beneficial for many machine learning algorithms. I'll display the first few rows of the scaled data afterwards.

The numerical features (age, annual_income, loan_amount, credit_score, num_dependents, and existing_loans_count) have been successfully scaled using `StandardScaler`. This transformation normalizes the data to have a mean of 0 and a standard deviation of 1, which is crucial for many machine learning algorithms. The output above shows the first few rows of the DataFrame with these scaled numerical features.

Split the dataset into training and testing sets.

Okay, I'll split the dataset into training and testing sets. First, I'll separate the features from the target variable (`loan_approved`). Then, I'll apply one-hot encoding to any categorical features to convert them into a numerical format suitable for machine learning models. Finally, I'll use `train_test_split` to create the training and testing portions of your data, and I'll show you the shape of each resulting set.

The dataset has been successfully split into training and testing sets. `X_train` and `X_test` contain the features, while `y_train` and `y_test` contain the target variable (`loan_approved`).

- The training set (`X_train` , `y_train`) consists of 800 samples.
- The testing set (`X_test` , `y_test`) consists of 200 samples.

`X_train` now includes one-hot encoded categorical features, as shown in its head, and `y_train` shows the corresponding loan approval status for these training samples.

📄 🔄

Train a Logistic Regression model using `X_train` and `y_train`

Evaluate the model performance using a confusion matrix

Visualize the feature importance from the trained model

What can I help you build?

+ 🌐

Gemini 2.5 Flash ▾ ▶

Gemini can make mistakes so double-check it and use code with caution. [Learn more](#)

Task 2 – Social Media Profiles

Preprocess a social media engagement dataset using AI assistance.

Instructions:


- Remove duplicate user profiles.
- Handle missing values
- Encode categorical columns

Dataset link:

<https://www.kaggle.com/datasets/medaxone/top-100-social-media-profiles>

Expected Output:

- A cleaned dataset for engagement prediction.



The image shows a Jupyter Notebook on the left and the Gemini AI chat interface on the right. The notebook contains Python code for preprocessing a dataset. The code includes steps for loading the dataset, cleaning column names, checking for missing values, imputing missing values with the median, cleaning the 'PROFILE' column, and cleaning the 'Platform' column. The Gemini chat interface on the right shows a conversation where the user asks to preprocess the 'SocialMediaTop100' dataset. The Gemini AI responds with a detailed explanation of the preprocessing steps, including cleaning column names, handling missing values, and cleaning the 'PROFILE' and 'Platform' columns. The AI also provides a summary of the actions performed and a list of the cleaned dataset's features.

```
import pandas as pd
import re

# Load the dataset
data = pd.read_csv('/content/drive/MyDrive/SocialMediaTop100.csv', encoding='latin1')

print('\nOriginal DataFrame Info:')
data.info()

# Clean column names by stripping whitespace
data.columns = data.columns.str.strip()

# Display cleaned column names
print('\nCleaned Column Names:')
print(data.columns)

# Check for missing values
print('\nMissing values before cleaning:')
print(data.isnull().sum())

# Impute missing values for 'FOLLOWERS' and 'POSTS' with the median
for col in ['FOLLOWERS', 'POSTS']:
    if data[col].isnull().any():
        median_val = data[col].median()
        data[col] = data[col].fillna(median_val)

print('\nMissing values after imputation:')
print(data.isnull().sum())

# Clean 'PROFILE' column: remove special unicode characters and strip whitespace
data['PROFILE'] = data['PROFILE'].apply(lambda x: re.sub(r'[\t\r\n]+', '', x).strip() if isinstance(x, str) else x)

# Clean 'Platform' column: strip whitespace
data['Platform'] = data['Platform'].apply(lambda x: x.strip() if isinstance(x, str) else x)

# Convert 'FOLLOWERS' and 'POSTS' to integer type after imputation
data['FOLLOWERS'] = data['FOLLOWERS'].astype(int)
data['POSTS'] = data['POSTS'].astype(int)

print('\nDataFrame head after cleaning:')
display(data.head())

print('\nDataFrame info after cleaning:')
data.info()
```

Gemini

preprocess the SocialMediaTop100 dataset

I will preprocess the SocialMediaTop100 dataset by first cleaning column names, then checking data types and handling any missing values. After that, I'll clean the 'PROFILE' and 'Platform' columns to remove special characters and strip extra whitespace.

It looks like the preprocessing steps failed because the `data` DataFrame was not loaded into the environment. I'll combine the data loading and preprocessing steps into a single code cell to make sure everything runs in the correct order.

It seems there was an issue applying the previous changes. The preprocessing steps are still encountering a `NameError` because the `data` DataFrame is not defined. I will re-consolidate all the data loading and preprocessing code into a single cell (`f4c7u2v0Jnv`) and clear the other problematic cells to fix this.

It looks like we're still encountering a `NameError` because the `data` DataFrame isn't recognized when the preprocessing steps are run. It seems my previous attempt to consolidate the code didn't fully take effect. I will re-apply the changes, moving all data loading and preprocessing into a single cell and clearing the others, to ensure the code executes correctly.

Okay, I've consolidated all the preprocessing steps for the `SocialMediaTop100` dataset into a single cell. I'm now executing that cell to clean the column names, handle missing values, remove special characters from the 'PROFILE' column, strip whitespace from the 'Platform' column, and convert 'FOLLOWERS' and 'POSTS' to integer types. Once executed, I will display the cleaned DataFrame's head and info.

The `SocialMediaTop100` dataset has been successfully preprocessed. I've performed the following actions:

- **Column Name Cleaning:** All column names were stripped of leading/trailing whitespace.
- **Missing Value Imputation:** The `POSTS` column, which initially had 100 missing values, had them imputed using the median value. The `FOLLOWERS` column had no missing values.

What can I help you build?

+ Gemini 2.5 Flash

Gemini can make mistakes so double-check it and use code with caution. [Learn more](#)

Original DataFrame Info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 500 entries, 0 to 499

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	N	500 non-null	int64
1	PROFILE	500 non-null	object
2	FOLLOWERS	500 non-null	float64
3	POSTS	400 non-null	float64
4	Platform	500 non-null	object

dtypes: float64(2), int64(1), object(2)

memory usage: 19.7+ KB

Cleaned Column Names:

```
Index(['N', 'PROFILE', 'FOLLOWERS', 'POSTS', 'Platform'], dtype='object')
```

Missing values before cleaning:

N	0
PROFILE	0
FOLLOWERS	0
POSTS	100
Platform	0

dtype: int64

Missing values after imputation:

N	0
PROFILE	0
FOLLOWERS	0
POSTS	0
Platform	0

dtype: int64

DataFrame head after cleaning:

	N	PROFILE	FOLLOWERS	POSTS	Platform
0	1	Instagram	539446645	7202	Instagram
1	2	Cristiano Ronaldo	473864939	3338	Instagram
2	3	Kylie	364542529	6935	Instagram
3	4	Leo Messi	355790796	890	Instagram
4	5	Selena Gomez	341579063	1828	Instagram

DataFrame info after cleaning:

```
<class 'pandas.core.frame.DataFrame'>
```

Task 3 – Weather Dataset Feature Engineering

Task:

Use AI to preprocess a weather dataset for forecasting.

Instructions:

- Convert date column into datetime format.
- Handle missing values in temperature, humidity.
- Create new features (season, week_number).
- Normalize rainfall values.

Dataset link:

<https://www.kaggle.com/code/alsamaualalhinai/weather-dataset/input>

Expected Output:

- A processed dataset ready for forecasting models.

The screenshot displays a Jupyter Notebook on the left and the Gemini AI interface on the right. The notebook contains Python code for preprocessing a weather dataset, including handling missing values and encoding categorical variables. The output of the code is visible in the notebook cells, showing the head of the DataFrame and the value counts for the 'Rain_encoded' column.

Final DataFrame head after preprocessing:

Temperature	Humidity	Wind_Speed	Cloud_Cover	Pressure	Rain	Rain_encoded
0	0	0	0	0	rain	1
1	0	0	0	0	no rain	0
2	0	0	0	0	no rain	0
3	0	0	0	0	rain	1
4	0	0	0	0	no rain	0

Value counts for 'Rain_encoded':

Rain_encoded	count
0	2186
1	314

The Gemini AI interface on the right shows a successful completion message and a summary of the preprocessing steps performed, including handling missing values, encoding categorical variables, and creating new features.

Task 4 – Movie Ratings Dataset Preparation

Task:

Clean and preprocess a movie ratings dataset using AI-generated scripts.

Instructions:

- Handle missing values in rating, genre.

- Encode categorical variables like genre, language.
- Remove duplicate movie entries.
- Scale rating values between 0 and 1.

Dataset link:

<https://www.kaggle.com/datasets/PromptCloudHQ/imdb-data>

Expected Output:

- A dataset suitable for recommendation systems.

```
# Display basic information about the DataFrame to identify missing values and data types
print("Original DataFrame Info:")
data.info()

# Check for missing values
print("\nMissing values before imputation:")
print(data.isnull().sum())
```

... Original DataFrame Info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Rank                  1000 non-null  int64
 1   Title                 1000 non-null  object
 2   Genre                 1000 non-null  object
 3   Description            1000 non-null  object
 4   Director              1000 non-null  object
 5   Actors                1000 non-null  object
 6   Year                  1000 non-null  int64
 7   Runtime (Minutes)     1000 non-null  int64
 8   Rating                1000 non-null  float64
 9   Votes                 1000 non-null  int64
10   Revenue (Millions)    872 non-null   float64
11   Metascore             936 non-null   float64
dtypes: float64(3), int64(4), object(5)
memory usage: 93.9+ KB
```

Missing values before imputation:

```
Rank          0
Title         0
Genre         0
Description   0
Director      0
Actors        0
Year          0
Runtime (Minutes)  0
Rating        0
Votes         0
Revenue (Millions)  128
Metascore     64
dtype: int64
```

```
# Impute missing 'Rating' values with the median
data['Rating'] = data['Rating'].fillna(data['Rating'].median())

# Impute missing 'Genre' values with 'Unknown'
data['Genre'] = data['Genre'].fillna('Unknown')

# Check for missing values after imputation
print("\nMissing values after imputation:")
print(data[['Rating', 'Genre']].isnull().sum())
```

```
Missing values after imputation:
Rating    0
Genre     0
dtype: int64
```

```
# Remove duplicate movie entries
duplicates_before = data.duplicated().sum()
data.drop_duplicates(inplace=True)
duplicates_after = data.duplicated().sum()

print(f"Number of duplicate entries before removal: {duplicates_before}")
print(f"Number of duplicate entries after removal: {duplicates_after}")
```

```
Number of duplicate entries before removal: 0
Number of duplicate entries after removal: 0
```

```
from sklearn.preprocessing import MinMaxScaler

# Scale 'Rating' values between 0 and 1
scaler = MinMaxScaler()
data['Rating_scaled'] = scaler.fit_transform(data[['Rating']])

print("\nHead of DataFrame with scaled 'Rating_scaled' column:")
display(data[['Rating', 'Rating_scaled']].head())
```

Head of DataFrame with scaled 'Rating_scaled' column:

	Rating	Rating_scaled
0	8.1	0.873239
1	7.0	0.718310
2	7.3	0.760563
3	7.2	0.746479
4	6.2	0.605634

he 'Genre' column contains multiple genres separated by commas. To encode this, we'll first split the genres and then apply one-hot encoding.

```
import numpy as np

# Split 'Genre' column into individual genres
data['genre'] = data['Genre'].apply(lambda x: x.split(',') if isinstance(x, str) else [x])

# Create a list of all unique genres
all_genres = np.unique([genre for sublist in data['genre'] for genre in sublist])

# Create one-hot encoded columns for each genre
for genre in all_genres:
    data[genre] = data['genre'].apply(lambda genres: 1 if genre in genres else 0)

# Drop the original 'Genre' column if desired
data.drop('Genre', axis=1, inplace=True)

print("\nDataFrame head after one-hot encoding 'Genre' and scaling 'Rating':")
display(data.head())
```

DataFrame head after one-hot encoding 'Genre' and scaling 'Rating':

Rank	Title	Description	Director	Actors	Year	Runtime (Minutes)	Rating	Votes	Revenue (millions)	...	Horror	Music	Musical	Mystery	Romance	Sci-Fi
0	1	Guardians of the Galaxy	A group of intergalactic criminals are forced...	James Gunn	Chris Pratt, Vin Diesel, Bradley Cooper, Zoe S...	2014	121	8.1	757074	333.13	...	0	0	0	0	1
1	2	Prometheus	Following clues to the origin of mankind, a te...	Ridley Scott	Noomi Rapace, Logan Marshall-Green, Michael Fa...	2012	124	7.0	485820	126.46	...	0	0	0	1	1
2	3	Split	Three girls are kidnapped by a man with a diag...	M. Night Shyamalan	James McAvoy, Anya Taylor-Joy, Haley Lu Richar...	2016	117	7.3	157606	135.12	...	1	0	0	0	0
3	4	Sing	In a city of humanoid animals, a hustling thea...	Christopher Lord	Matthew McConaughey, Reese Witherspoon, Seth Ma...	2016	108	7.2	69545	270.32	...	0	0	0	0	0
4	5	Suicide Squad	A secret government agency recruits some of th...	David Ayer	Will Smith, Jared Leto, Margot Robbie, Viola D...	2016	123	6.2	393727	325.02	...	0	0	0	0	0

5 rows x 32 columns

Summary

• Handle missing values in rating, genre. • Encode categorical variables like genre, language. • Remove duplicate movie entries. • Scale rating values between 0 and 1. for IMDB-Movie-Data.

• I will preprocess the IMDB-Movie-Data.csv dataset. This involves handling missing values in the 'Rating' and 'Genre' columns, removing any duplicate movie entries, and encoding the 'Genre' column. Finally, I will scale the 'Rating' values to a range between 0 and 1.

The preprocessing for the IMDB-Movie-Data.csv dataset has been successfully completed:

- **Missing Values:** I first checked the DataFrame and observed that 'Revenue' (Millions) and 'Metascore' had missing values, while 'Rating' and 'Genre' had none at the start. I then imputed any potential missing values in 'Rating' with its median and 'Genre' with 'Unknown' (though they were already complete).
- **Duplicate Entries:** I checked for and removed duplicate movie entries. It was found that there were no duplicate entries in the dataset.
- **Rating Scaling:** The 'Rating' values were successfully scaled to a range between 0 and 1, creating a new column named 'Rating_scaled'.
- **Genre Encoding:** The 'Genre' column was one-hot encoded, effectively creating new binary columns for each unique genre present in the dataset. The original 'Genre' column was then dropped.

The DataFrame is now preprocessed as requested, with all these transformations applied.

Excellent! The missing values in both the 'Rating' and 'Genre' columns have been successfully handled, and there are now no missing entries in either of these columns. This completes another step in the preprocessing of the IMDB-Movie-Data.csv dataset.

All preprocessing tasks for the loanapproval.csv, SocialMediaTop100.csv, and IMDB-Movie-Data.csv datasets have been successfully completed. If you have any new instructions or further analyses you'd like to perform, please let me know!

All preprocessing tasks for the loanapproval.csv, SocialMediaTop100.csv, weather_forecast_data.csv, and IMDB-Movie-Data.csv datasets have been successfully completed as instructed.

- For loanapproval.csv, missing values were imputed, 'loan_status' was encoded, categorical features were one-hot encoded, and numerical features were scaled.
- For SocialMediaTop100.csv, column names were cleaned, missing 'FOLLOWERS' and 'POSTS' were imputed, and 'PROFILE' and 'Platform' columns were cleaned.
- For weather_forecast_data.csv, missing 'Temperature' and 'Humidity' values were checked, and the 'Rain' column was numerically encoded.
- For IMDB-Movie-Data.csv, missing 'Rating' and 'Genre' values were handled, duplicates were checked, 'Rating' was scaled, and 'Genre' was one-hot encoded.

All tasks requested in this session have been completed. If you have any new analyses or further questions, please let me know!

What can I help you build?

+ 55

Gemini 2.5 Flash

Gemini can make mistakes so double-check it and use code with caution. [Learn more](#)